

Deep Learning lecture 9 Sequence Modeling (2)

Yi Wu, IIIS

Spring 2025


Apr-14

Today's Topic

- Sequence to Sequence Model and Attention Mechanism
- The Transformer Model
- Generation Speedup for Transformer Model



RNN Recap

- Recurrent Neural Network
 - Same MLP network over a sequence (i.e., “loops”)
 - Arbitrarily long sequences \rightarrow fixed-sized vector
 - Training: backpropagation through time (BPTT)
 - Practical Issues
 - Weights/Gradient explosion and saturation
 - A few tricks for gradient explosion
 - Gradient clipping, truncated BPTT, careful initialization
- Long Short-Term Memory (LSTM) Network
 - A specialized RNN for long-term dependency (~ 100 timesteps)
 - Key ideas: elementary gates
 - Variants: bidirectional LSTM; Peephole LSTM; GRU; etc



RNN Recap

- Autoregressive Language Model
 - Generative model over texts: $P(X) = \prod_t P(X_t | X_{i < t})$
 - LSTM language model: $Y_t, h_t = LSTM(h_{t-1}, X_t)$; $P(X_t | X_{i < t}) = \text{Softmax}(Y_t)$
 - Word Embedding
 - A distributed representation for word semantics
 - Word2Vec: a tool for word embedding
 - Objective: from context c to predict word w
 - CBOW and Skip-Gram
 - Negative Sampling
 - Multi-class prediction \rightarrow binary classification
 - D training corpus; V vocabulary
- $$L(W, C) = \sum_{(c, w) \in D} \log \frac{1}{\exp(-w^T c) + 1} + \sum_{c \in D, w \in V} \log \frac{\exp(-\tilde{w}^T c)}{\exp(-\tilde{w}^T c) + 1}$$



RNN Recap

- Autoregressive Language Model



- Generative model over texts: $P(X) = \prod_t P(X_t | X_{i < t})$
 - LSTM language model: $Y_t, h_t = LSTM(h_{t-1}, X_t)$; $P(X_t | X_{i < t}) = \text{Softmax}(Y_t)$
- Word Embedding
 - A distributed representation for word semantics

- Word2Vec: a tool for word embedding

- Objective: from context c to predict word w
 - CBOW and Skip-Gram
 - Negative Sampling
 - Multi-class prediction \rightarrow binary classification
 - D training corpus; V vocabulary

$$L(W, C) = \sum_{(c, w) \in D} \log \frac{1}{\exp(-w^T c) + 1} + \sum_{c \in D, w \in V} \log \frac{\exp(-\tilde{w}^T c)}{\exp(-\tilde{w}^T c) + 1}$$


Copyright © IIIS, Tsinghua University



Word2Vec is only for representation learning!
It does not care about prediction accuracy!

RNN Recap

- Autoregressive Language Model
 - Generative model over texts: $P(X) = \prod_t P(X_t | X_{i < t})$
 - LSTM language model: $Y_t, h_t = LSTM(h_{t-1}, X_t)$; $P(X_t | X_{i < t}) = \text{Softmax}(Y_t)$
 - Word Embedding
 - A distributed representation for word semantics
- Word2Vec: a tool for word embedding
- More Techniques
 - Hierarchical softmax
 - Beam search
 - ELMo for contextualized embeddings



Language Model Applications

- Text Classification
 - Supervised learning
- Text Generation
 - $p(X; \theta)$: the probability for X
 - Unconditioned Generation
 - E.g., AI作诗
 - **Conditioned generation?**
 - E.g., Machine translation

深度之梦

在数据的海洋里遨游，
算法如风，吹散迷雾。
神经元闪烁似星辰，
连接着未来的道路。
梯度回溯千重浪，
优化求解万象生。
一行代码塑乾坤，
模型自我去提升。



Deep learning is a popular area in AI.

检测语言 英语 中文 德语

中文 (简体) 英语 日语

深度学习是AI的热门领域。

Shèndù xuéxí shì AI de rèmén lǐngyù.

4/13 38 / 500 Copyright © IIIS, Tsinghua University 7

Machine Translation

- A task of translating a sentence from a source language to the target language

x: *L'homme est né libre, et partout il est dans les fers*

y: *Man is born free, but everywhere he is in chains*

人生而自由，却无往不在枷锁之中。——卢梭《社会契约论》


– Rousseau

Machine Translation

- Before 2014: Statistical Machine Translation
 - Extremely complex systems that require massive human efforts
 - Separately designed components
 - A lot of feature engineering
 - Lots of linguistics domain knowledge and expertise
- Before 2016:
 - Google's commercial translation product is based on statistical machine translation
- What happened in 2014?
 - A borrowed slide from Stanford CS224

2014

(dramatic reenactment)



2014

MT research
Learning, Tsinghua U

Copyright © IIIS, Tsinghua University

Neural
Machine
Translation

(dramatic reenactment)

Sequence to Sequence Model


- Neural Machine Translation (NMT)
 - Learning to translate via single end-to-end neural network!
 - Source language X , then $Y = f(X; \theta)$
- Sequence-to-Sequence Model (Seq2Seq, Sutskever et al, NIPS2014)
 - NeurIPS 2024 test-of-time award
 - Two RNNs: f_{enc} and f_{dec} , $X \rightarrow f_{enc} \rightarrow h \rightarrow f_{dec} \rightarrow Y$
 - Encoder f_{enc}
 - It takes in X , and produce the initial hidden state h for decoder
 - We can use bidirectional RNN
 - Decoder f_{dec}
 - It takes in the hidden state h from f_{enc} to generate Y
 - Autoregressive language model

Sequence to Sequence Model

The sequence-to-sequence model

Encoding of the source sentence.
Provides initial hidden state
for Decoder RNN.


Encoder RNN



Source sentence (input)

Encoder RNN produces
an **encoding** of the
source sentence.

Target sentence (output)



Decoder RNN

Decoder RNN is a Language Model that generates target sentence, *conditioned on encoding*.


Note: This diagram shows **test time** behavior: decoder output is fed in as next step's input

Sequence to Sequence Model

- Seq2Seq is a conditioned language model
 - $h = f_{enc}(X)$ (final hidden state)
 - $Y = f_{dec}(h)$ (a LM that conditions on the initial hidden state h)
- Seq2Seq model is particularly generic for a lot of applications
 - Summarization (摘要) or Captioning (起标题)
 - Article \rightarrow abstract/caption
 - Dialogue (对话)
 - Previous utterance \rightarrow next utterance
 - Code generation
 - Natural language \rightarrow python
 - VAE-based seq2seq model for text generation with latent variables


Sequence to Sequence Model

- How to train a seq2seq model?
 - Collect a huge paired dataset and train it end-to-end via BPTT!
 - MLE learning for $P(Y|X) = P(Y|f_{enc}(X))$




Sequence to Sequence Model

- We can also make the model deeper!
 - Stacked seq2seq model






Sequence to Sequence Model

- 2016: Google switch google translate from SMT to NMT
 - Seq2Seq paper has >28.6k citations since 2014




Sequence to Sequence Model

- Issue in the vanilla Seq2Seq model
 - Alignment: the word-level correspondence between X and Y
 - There are complex long-term dependencies



Sequence to Sequence Model


- Issue in the vanilla Seq2Seq model
 - The information bottleneck due to h
 - We want each Y_t to also focus on X_i that it is aligned with




Sequence-to-Sequence with Attention

- NMT by Jointly Learning to Align and Translate
 - Bahdanau, Cho & Bengio, ICLR 2015 (38.5k citation)
 - Core idea:
 - When decoding Y_t , we consider both hidden states and alignments
 - Hidden: h_{t-1} from $Y_{i < t}$, i.e., $h_{t-1} = f_{dec}(Y_{i < t})$
 - Alignment: a direction connection to “key” words from X
 - Which part of X to focus?
 - Learn a softmax weight over X (attention distribution P_{att})
 - $P_{att}(X_i | h_{t-1})$: how much attention you want to put on word X_i
 - attention output $h_{att} = \sum_i f_{enc}(X_i | X_{j < i}) \cdot P_{att}(X_i | h_{t-1})$
 - Use h_{t-1} and h_{att} to compute Y_t
 - Let’s go through the diagram before showing more details


Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention



Sequence-to-Sequence with Attention



Sequence-to-Sequence with Attention




Sequence-to-Sequence with Attention

- Attention in equations
 - Input sequence X and encoder f_{enc} and decoder f_{dec}
 - $f_{enc}(X)$ produces hidden states $h_1^{enc}, h_2^{enc}, \dots, h_N^{enc}$
 - On timestep t , we have decoder hidden state h_t
 - Attention score $e_i = h_t^T h_i^{enc}$
 - Attention distribution $\alpha_i = P_{att}(X_i) = \text{softmax}(e_i)$
 - Attention output
- $h_{att}^{enc} = \sum_i \alpha_i h_i^{enc}$
- $Y_t \sim g(h_t, h_{att}^{enc}; \theta)$
 - Sample output using both h_t and h_{att}^{enc}

Attention

- Attention is great!
 - It significantly improves NMT!
 - It solves the bottleneck problem and long-term dependency issue
 - Also helps gradient vanishing problem
 - It provides some interpretability
 - We can understand the focus of RNN decoder
- Attention is a general technique
 - Given a set of vector values V_i and a vector query q
 - Attention computes a weighted sum of values depending on q
- Attention can learn a representation of an arbitrary set of vectors $\{v_i\}$ depending on query q



Attention



- Attention can learn a representation of an arbitrary set of vectors $\{v_i\}$ depending on query q
 - $\alpha_i = \text{softmax}(f(v_i, q))$ (attention distribution)
 - $v_{att} = \sum_i \alpha_i v_i$ (attention output)
 - Attention is *size-invariant* and *order-invariant*
- More use cases
 - E.g., a representation of a set of points (Pointer network, NIPS2015 & Deep Sets, NIPS2017)
 - E.g., include non-local information in CNN (Non-local network, CVPR18; Self-Attention GAN, ICML 19; BigGAN, ICLR 19)

Attention

- Attention can learn a representation of an arbitrary set of vectors $\{v_i\}$ depending on query q
 - $\alpha_i = \text{softmax}(f(v_i, q))$ (attention distribution)
 - $v_{att} = \sum_i \alpha_i v_i$ (attention output)
- Attention Variants $f(v_i, q)$
 - Multiplicative attention: $f(v_i, q) = q^T W h_i$
 - W is a weight matrix
 - Additive attention: $f(v_i, q) = u^T \tanh(W_1 v_i + W_2 q)$
 - W_1, W_2 are weight matrices
 - u is a weight vector
 - Expressiveness v.s. efficiency


Sequence to Sequence Model with Attention

- Attention-Based Seq2Seq Model
 - Use bidirectional LSTMs as your encoder for input data
 - Use stacked LSTMs as your decoder for output data
 - Use attention for long-term dependencies



Sequence to Sequence Model with Attention

- Attention-Based Seq2Seq Model
 - Use bidirectional LSTMs as your encoder for input data
 - Use stacked LSTMs as your decoder for output data
 - Use attention for long-term dependencies
- Most NLP applications!
 - till 2017 😊



2014-2017ish

Recurrence


Lots of trial
and error



?????


Sequence to Sequence Model with Attention

- Story So Far
 - RNN Models
 - Simple & Generic solution for sequence modeling
 - Issue for long-term dependencies
 - Linear computations for distant words through a single latent state
 - Lack of parallelization
 - Forced sequential computation (contrast with CNN)




Sequence to Sequence Model with Attention

- Story So Far
 - RNN Models
 - Simple & Generic solution for sequence modeling
 - Issue for long-term dependencies
 - Linear computations for distant words through a single latent state
 - Lack of parallelization
 - Forced sequential computation (contrast with CNN)
 - Attention
 - Direct connection to distant words
 - $O(N)$ computation but perfectly parallel!
 - **Attention is all we need?**



Self-Attention

- Attention is all you need (NIPS2017, Google Brain)
 - A purely attention-base architecture for sequence modeling
 - NO RNN at all
 - Basic component: Self-Attention, $Y = f_{SA}(X; \theta)$
 - Core idea:
 - X_t attends on the entire X sequence
 - Y_t computed from X_t and the attention output
 - Equations for Y_t
 - Key k_t , value v_t , query q_t from X_t
 - $k_t, v_t, q_t = g_1(X_t; \theta)$
 - Attention distribution $\alpha_{t,j} = \text{softmax}(q_t^T k_j)$
 - Attention output $out_t = \sum_j \alpha_{t,j} v_j$
 - $Y_t = g_2(out_t; \theta)$
- Issues of self-attention?



Self-Attention

- Issues of Vanilla Self-Attention
 - Notion of sequence order
 - Attention is order-invariant
 - Lack of non-linearities
 - All the weights are simple linear weighted average
 - Capability of autoregressive modelling
 - In generation tasks, the model cannot “look at the future”
 - E.g., text generation
 - Y_t can only depend on $X_{i < t}$
 - Vanilla self-attention focuses on the entire sequence

Self-Attention

- Issues of Vanilla Self-Attention
 - **Notion of sequence order**
 - Attention is order-invariant
 - Lack of non-linearities
 - All the weights are simple linear weighted average
 - Capability of autoregressive modelling
 - In generation tasks, the model cannot “look at the future”
 - E.g., text generation
 - Y_t can only depend on $X_{i < t}$
 - Vanilla self-attention focuses on the entire sequence



Self-Attention

- Notion of Sequence Ordering
 - Vanilla attention
 - $\tilde{\alpha}_{i,j} = \text{softmax}(\tilde{q}_i^T \tilde{k}_j)$; $\text{out}_i = \sum_j \tilde{\alpha}_{i,j} \tilde{v}_j$
 - $\tilde{k}_t, \tilde{v}_t, \tilde{q}_t = g_1(X_t)$, do not contain position information
 - Idea: position encoding
 - p_i : an embedding vector of position i
 - $k_t, v_t, q_t = g_1([X_t, p_t])$ include position features
 - Practical remark:
 - Additive can be sufficient: $k_t \leftarrow \tilde{k}_t + p_t, q_t \leftarrow \tilde{q}_t + p_t, v_t \leftarrow \tilde{v}_t + p_t$
 - p_t is typically only included in the first layer
 - **How to design p_i ?**
 - Note that the length of a sequence can be long

Self-Attention

- Notion of Sequence Ordering
 - Idea: position encoding p_i
 - $\tilde{k}_t, \tilde{v}_t, \tilde{q}_t = g_1(X_t)$, do not contain position information
 - Design of p_t
 - **Sinusoidal position representation**
 - Concatenate sinusoidal functions of varying periods

$$p_i = \begin{pmatrix} \sin(i/10000^{2*1/d}) \\ \cos(i/10000^{2*1/d}) \\ \vdots \\ \sin(i/10000^{2*\frac{d}{2}/d}) \\ \cos(i/10000^{2*\frac{d}{2}/d}) \end{pmatrix}$$

Heatmap of $p_i^T p_j$ 

OpenPsi @ IIIS

Self-Attention

- Notion of Sequence Ordering
 - Idea: position encoding p_i
 - $\tilde{k}_t, \tilde{v}_t, \tilde{q}_t = g_1(X_t)$, do not contain position information
 - Design of p_t
 - Sinusoidal position representation
 - **Learned absolute representation**
 - Let p_t become a learned parameter vector!
 - Assume maximum length L , learn a matrix $p \in \mathbb{R}^{d \times T}$, p_t is a column of p
 - A popular choice in practice!
 - Pros:
 - Flexible and learnable, more powerful
 - Cons:
 - Assume a fixed maximum length L , does not work at all for length above L

Self-Attention


- Notion of Sequence Ordering
 - Idea: position encoding p_t
 - $\tilde{k}_t, \tilde{v}_t, \tilde{q}_t = g_1(X_t)$, do not contain position information
 - Design of p_t
 - Sinusoidal position representation
 - Learned absolute representation
 - **Relative position representation (ACL2018, Google)**
 - When computing attention, relative distance is important!
 - $\alpha_{i,j} = \text{softmax}(q_i^T(k_j + p_{[i-j]}))$
 - $out_i = \sum_j \alpha_{i,j} (v_j + p_{[i-j]})$
 - Bounded relative distance $p_t = p_{\max(-k, \min(k, t))}$
 - Truncate $t < -k$ to k and $t > k$ to k
 - Pros: learned representation and extrapolate well; More powerful.
 - Cons: computation overhead (refer to the paper for implementation tricks)

Self-Attention

- Notion of Sequence Ordering
 - Idea: position encoding p_i
 - $\tilde{k}_t, \tilde{v}_t, \tilde{q}_t = g_1(X_t)$, do not contain p_i
 - Design of p_t
 - Sinusoidal position representation
 - Learned absolute representation
 - Relative position representation
 - Rotary position embedding (RoPE, Roformer, 2021)**
 - Relative position but factored computation

$$R_{\Theta,m}^d = \begin{pmatrix} \cos m\theta_1 & -\sin m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ \sin m\theta_1 & \cos m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & \cos m\theta_2 & -\sin m\theta_2 & \cdots & 0 & 0 \\ 0 & 0 & \sin m\theta_2 & \cos m\theta_2 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & \cos m\theta_{d/2} & -\sin m\theta_{d/2} \\ 0 & 0 & 0 & 0 & \cdots & \sin m\theta_{d/2} & \cos m\theta_{d/2} \end{pmatrix}$$

Copyright © IIIS, Tsinghua University



$$\Theta = \{\theta_i = 10000^{-2(i-1)/d}, i \in [1, 2, \dots, d/2]\}$$

$$q_m^T k_n = (R_{\Theta,m}^d W_q x_m)^T (R_{\Theta,n}^d W_k x_n) = x^T W_q R_{\Theta,n-m}^d W_k x_n$$

$$R_{\Theta,n-m}^d = (R_{\Theta,m}^d)^T R_{\Theta,n}^d$$

Remark:

- Compatible with any dimension and length
- Fast computation
- Effective in practice

Self-Attention


- Issues of Vanilla Self-Attention
 - **Notion of sequence order**
 - Solution: position encoding
 - Lack of non-linearities
 - All the weights are simple linear weighted average
 - Capability of autoregressive modelling
 - In generation tasks, the model cannot “look at the future”
 - E.g., text generation
 - Y_t can only depend on $X_{i < t}$
 - Vanilla self-attention focuses on the entire sequence

Self-Attention

- Issues of Vanilla Self-Attention
 - Notion of sequence order
 - Solution: position encoding
 - **Lack of non-linearities**
 - All the weights are simple linear weighted average
 - Capability of autoregressive modelling
 - In generation tasks, the model cannot “look at the future”
 - E.g., text generation
 - Y_t can only depend on $X_{i < t}$
 - Vanilla self-attention focuses on the entire sequence

Self-Attention

- Combine nonlinearities in self-attention
 - Vanilla self-attention
 - No element-wise activation functions (e.g., ReLU, tanh)
 - Only weighted average and softmax operators
 - Essentially linear transformations of inputs
 - Easy fix:
 - Add an MLP to process out_i
 - $m_i = MLP(out_i)$
 - $= W_2 \cdot \text{ReLU}(W_1 \cdot out_i + b_1) + b_2$
 - Remark
 - we do not put activation layer before softmax



Self-Attention


- Issues of Vanilla Self-Attention
 - Notion of sequence order
 - Solution: position encoding
 - **Lack of non-linearities**
 - Solution: post-processing MLP layer
 - Capability of autoregressive modelling
 - In generation tasks, the model cannot “look at the future”
 - E.g., text generation
 - Y_t can only depend on $X_{i < t}$
 - Vanilla self-attention focuses on the entire sequence

Self-Attention

- Issues of Vanilla Self-Attention
 - Notion of sequence order
 - Solution: position encoding
 - Lack of non-linearities
 - Solution: post-processing MLP layer
 - **Capability of autoregressive modelling**
 - In generation tasks, the model cannot “look at the future”
 - E.g., text generation
 - Y_t can only depend on $X_{i < t}$
 - Vanilla self-attention focuses on the entire sequence

Self-Attention

- Autoregressive Modeling
 - In language mode decoder: model $P(Y_t | X_{i < t})$
 - out_t cannot look at future $X_{i > t}$
 - Naïve solution:
 - For each t , a varying for-loop only iterating over $i \leq t$
 - Varying for-loop for each t , parallelization unfriendly
 - Masked Attention
 - Compute $e_{i,j} = q_i^T k_j$ as usual (perfect parallel)
 - Mask out $e_{i>j}$ by setting $e_{i>j} = -\infty$ (perfect parallel)
 - $e \odot (1 - M) \leftarrow -\infty$; M is a fixed 0/1 mask matrix
 - Then compute $\alpha_i = \text{softmax}(e_i)$ (perfect parallel)
 - Remark:
 - $M = 1$ for full self-attention
 - Set M for arbitrary dependency ordering



Self-Attention

- Issues of Vanilla Self-Attention


- Notion of sequence order
 - Solution: position encoding
- Lack of non-linearities
 - Solution: post-processing MLP layer
- **Capability of autoregressive modelling**
 - Solution: masked self-attention

Self-Attention

- Issues of Vanilla Self-Attention
 - Notion of sequence order
 - Solution: position encoding
 - Lack of non-linearities
 - Solution: post-processing MLP layer
 - Capability of autoregressive modelling
 - Solution: masked self-attention
- Basic building block for the famous “Transformer” model!
 - Attention is all you need (NIPS2017, Vaswani et al, Google)
 - Self-attention + **a few more other enhancements!**
 - A milestone: first pure attention-based model for effective sequence modeling
 - *Originally proposed for NMT; Soon dominates general sequence modeling problems*

Transformer Model

- Transformer-based sequence-to-sequence modeling



Transformer Model

- Transformer-based sequence-to-sequence modeling
 - Basic building blocks: masked self-attention
 - Enhancements
 - **Key-query-value attention**
 - Obtain q_t, v_t, k_t from X_t
 - $q_t = W^q X_t; v_t = W^v X_t; k_t = W^k X_t$ (position encoding omitted)
 - W^q, W^v, W^k are learnable weight matrices
 - $\alpha_{i,j} = \text{softmax}(q_i^T k_j); out_i = \sum_j \alpha_{i,j} v_j$
 - Intuition: key, query, and value can focus on different parts of input

All pairs of attention scores!

$$\begin{aligned}
 XQ & \quad K^T X^T = XQK^T X^T \in \mathbb{R}^{T \times T} \\
 \text{softmax} \left(XQK^T X^T \right) & \quad XV = \text{output} \in \mathbb{R}^{T \times d}
 \end{aligned}$$

Transformer Model

- Transformer-based sequence-to-sequence modeling
 - Basic building blocks: masked self-attention
 - Enhancements
 - Key-query-value attention
 - **Multi-headed attention**
 - Standard attention → single-headed attention
 - $X_t \in \mathbb{R}^d, Q, K, V \in \mathbb{R}^{d \times d}$
 - We only “look at” a single position j with high $\alpha_{i,j}$
 - What if we want to look at different j for different reasons?
 - Idea: define h separate attention heads
 - h different attention distributions, keys, values and queries
 - $Q^l, K^l, V^l \in \mathbb{R}^{d \times \frac{d}{h}}$, for $1 \leq l \leq h$
 - $\alpha_{i,j}^l = \text{softmax}(q_i^l k_j^l); out_i^l = \sum_j \alpha_{i,j}^l v_j^l$

#Params Unchanged!

Single-head attention
(just the query matrix)

$$X \quad Q = XQ$$

Multi-head attention
(just two heads here)


$$X \quad Q_1 \quad Q_2 = XQ_1 \quad XQ_2$$

Transformer Model

- Transformer-based sequence-to-sequence
 - Basic building blocks: masked self-attention
 - Enhancements
 - Key-query-value attention
 - **Multi-headed attention**
 - Standard attention → single-headed attention
 - $X_t \in \mathbb{R}^d$, $Q, K, V \in \mathbb{R}^{d \times d}$
 - We only “look at” a single position j with high probability
 - What if we want to look at different j for different i ?
 - Idea: define h separate attention heads
 - h different attention distributions, keys, values
 - $Q^l, K^l, V^l \in \mathbb{R}^{d \times \frac{d}{h}}$, for $1 \leq l \leq h$
 - $\alpha_{i,j}^l = \text{softmax} \left(q_i^l k_j^l \right)$; $out_i^l = \sum_j \alpha_{i,j}^l v_j^l$

Utterance Level Representation

$$\underline{c} = [c_1 \ c_2 \ c_3]$$




Sequence of Encoded Representations or Hidden States

Transformer Model


- Transformer-based sequence-to-sequence modeling
 - Basic building blocks: masked self-attention
 - Enhancements
 - Key-query-value attention
 - Multi-headed attention
 - **Architecture modifications**
 - Residual connection
 - Layer normalization
 - $out_t = \text{LN}(f_{SA}(X_t, M) + X_t); m_t = \text{LN}(\text{MLP}(out_t) + out_t)$
 - Scaled dot product
 - Intuition: when dimension d becomes large, $q^T k$ can be large
 - Issue: input to softmax can be large and make gradient small
 - $\alpha_{i,j}^l = \text{softmax} \left(\frac{q_i^l k_j^l}{\sqrt{d/h}} \right)$

Transformer Model

- Transformer-based sequence-to-sequence modeling
 - Basic building blocks: masked self-attention
 - Position encoding
 - Post-processing MLP
 - Attention mask
 - Enhancements
 - Key-query-value attention
 - Multi-headed attention
 - Architecture modifications
 - Residual connection
 - Layer normalization
 - Scaled dot product




Transformer-Based Seq2Seq Model



Transformer-Based Seq2Seq Model

Cross-attention


Conditioning on encoder hidden states!



Transformer-Based Seq2Seq Model

- Cross-attention
 - The conditioning part of transformer
 - Decoder can generate texts conditioning on the input sequence
 - Just like standard attention in RNN seq2seq model
 - z_t : decoder SA module inputs
 - h_t : encoder output hidden states
 - For each decoder out_t , we attend on encoder hiddens
 - Query from decoder: $q_t = W^q z_t$
 - Key and value from encoder: $k_j = W^k h_j$; $v_j = W^v h_j$
 - $\alpha_{ij} = \text{softmax}\left(\frac{q_t^T k_j}{\sqrt{d}}\right)$; $out_i = LN\left(\sum_j \alpha_{ij} v_j + z_i\right)$
 - Many practical variants can be implemented

Transformer-Based Seq2Seq Model



Transformer-Based Seq2Seq Model

- Machine translation with transformer (NIPS2017, Google)

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1		$3.3 \cdot 10^{18}$
Transformer (big)	28.4	41.8		$2.3 \cdot 10^{19}$


Transformer-Based Seq2Seq Model

- Generating Wikipedia by summarizing long sequences (ICLR2018, Google)
 - Document generation

Model	Test perplexity	ROUGE-L
<i>seq2seq-attention, $L = 500$</i>	5.04952	12.7
<i>Transformer-ED, $L = 500$</i>	2.46645	34.2
<i>Transformer-D, $L = 4000$</i>	2.22216	33.6
<i>Transformer-DMCA, no MoE-layer, $L = 11000$</i>	2.05159	36.2
<i>Transformer-DMCA, MoE-128, $L = 11000$</i>	1.92871	37.9
<i>Transformer-DMCA, MoE-256, $L = 7500$</i>	1.90325	38.8

Transformer Model for Images

- Vision Transformer (ViT, Google Brain, ICLR 2021, 33.4k citation)
 - *An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale*
 - Decompose an image to 16x16 patches and then apply transformer encoder




Transformer Model for Images

- Vision Transformer (ViT, Google Brain, ICLR 2021, 33.4k citation)
 - *An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale*
 - Decompose an image to 16x16 patches and then apply transformer encoder

	Ours-JFT (ViT-H/14)	Ours-JFT (ViT-L/16)	Ours-I21k (ViT-L/16)	BiT-L (ResNet152x4)	Noisy Student (EfficientNet-L2)
ImageNet	88.55 \pm 0.04	87.76 \pm 0.03	85.30 \pm 0.02	87.54 \pm 0.02	88.4/88.5*
ImageNet ReaL	90.72 \pm 0.05	90.54 \pm 0.03	88.62 \pm 0.05	90.54	90.55
CIFAR-10	99.50 \pm 0.06	99.42 \pm 0.03	99.15 \pm 0.03	99.37 \pm 0.06	—
CIFAR-100	94.55 \pm 0.04	93.90 \pm 0.05	93.25 \pm 0.05	93.51 \pm 0.08	—
Oxford-IIIT Pets	97.56 \pm 0.03	97.32 \pm 0.11	94.67 \pm 0.15	96.62 \pm 0.23	—
Oxford Flowers-102	99.68 \pm 0.02	99.74 \pm 0.00	99.61 \pm 0.02	99.63 \pm 0.03	—
VTAB (19 tasks)	77.63 \pm 0.23	76.28 \pm 0.46	72.72 \pm 0.21	76.29 \pm 1.70	—
TPUv3-core-days	2.5k	0.68k	0.23k	9.9k	12.3k

Transformer Model for Images

- Swin Transformer (MSRA, CVPR 2021 best paper)
 - Build hierarchical feature maps at different resolution
 - Self-attention only within each block (linear computation for image size)
 - Shifted block partitions to encode information between blocks



Transformer Model for Images

- Swin Transformer (MSRA, CVPR 2021 best paper)
 - Build hierarchical feature maps at different resolution
 - Self-attention only within each block (linear computation for image size)
 - Shifted block partitions to encode information between blocks

Method	mini-val		test-dev		#param.	FLOPs
	AP ^{box}	AP ^{mask}	AP ^{box}	AP ^{mask}		
RepPointsV2* [12]	-	-	52.1	-	-	-
GCNet* [7]	51.8	44.7	52.3	45.4	-	1041G
RelationNet++* [13]	-	-	52.7	-	-	-
SpineNet-190 [21]	52.6	-	52.8	-	164M	1885G
ResNeSt-200* [78]	52.5	-	53.3	47.1	-	-
EfficientDet-D7 [59]	54.4	-	55.1	-	77M	410G
DetectoRS* [46]	-	-	55.7	48.5	-	-
YOLOv4 P7* [4]	-	-	55.8	-	-	-
Copy-paste [26]	55.9	47.2	56.0	47.4	185M	1440G
X101-64 (HTC++)	52.3	46.0	-	-	155M	1033G
Swin-B (HTC++)	56.4	49.1	-	-	160M	1043G
Swin-L (HTC++)	57.1	49.5	57.7	50.2	284M	1470G
Swin-L (HTC++)*	58.0	50.4	58.7	51.1	284M	-


Table 2. Results on COCO object detection and instance segmentation. [†]denotes that additional deconvolution layers are used to produce hierarchical feature maps. * indicates multi-scale testing.

Method	Backbone	val mIoU	test score	ADE20K	
				#param.	FLOPs
DANet [23]	ResNet-101	45.2	-	69M	1119G
DLab.v3+ [11]	ResNet-101	44.1	-	63M	1021G
ACNet [24]	ResNet-101	45.9	38.5	-	-
DNL [71]	ResNet-101	46.0	56.2	69M	1249G
OCRNet [73]	ResNet-101	45.3	56.0	56M	923G
UperNet [69]	ResNet-101	44.9	-	86M	1029G
OCRNet [73]	HRNet-w48	45.7	-	71M	664G
DLab.v3+ [11]	ResNeSt-101	46.9	55.1	66M	1051G
DLab.v3+ [11]	ResNeSt-200	48.4	-	88M	1381G
SETR [81]	T-Large [‡]	50.3	61.7	308M	-
UperNet	DeiT-S [†]	44.0	-	52M	1099G
UperNet	Swin-T	46.1	-	60M	945G
UperNet	Swin-S	49.3	-	81M	1038G
UperNet	Swin-B [‡]	51.6	-	121M	1841G
UperNet	Swin-L [‡]	53.5	62.8	234M	3230G

Table 3. Results of semantic segmentation on the ADE20K val and test set. [†] indicates additional deconvolution layers are used to produce hierarchical feature maps. [‡] indicates that the model is pre-trained on ImageNet-22K.

Comparisons


- CNN v.s. LSTM v.s. Transformer



Context length L	Temporal Convolution	RNN	Transformer
Layers			
Generation			
Inference			

Comparisons


- CNN v.s. LSTM v.s. Transformer



Context length L	Temporal Convolution	RNN	Transformer
Layers	$O(\log L)$	$O(1)$	$O(1)$
Generation			
Inference			

Comparisons


- CNN v.s. LSTM v.s. Transformer



Context length L	Temporal Convolution	RNN	Transformer
Layers	$O(\log L)$	$O(1)$	$O(1)$
Generation	$O(L \log L)$	$O(L)$	$O(L^2)$
Inference			

Comparisons


- CNN v.s. LSTM v.s. Transformer



Context length L	Temporal Convolution	RNN	Transformer
Layers	$O(\log L)$	$O(1)$	$O(1)$
Generation	$O(L \log L)$	$O(L)$	$O(L^2)$
Inference	$O(\log L)$	$O(L)$	$O(1)$

Comparisons

- CNN v.s. LSTM v.s. Transformer



Context length L	Temporal Convolution	RNN	Transformer
Layers	$O(\log L)$	$O(1)$	$O(1)$
Generation	$O(L \log L)$	$O(L)$	$O(L^2)$
Inference	$O(\log L)$	$O(L)$	$O(1)$

Quadratic!!

Can we speedup transformer generation?

Speed up Transformers

- Quadratic generation cost
 - $O(L^2)$ for length L : sequential generation $O(L) \times$ attention $O(L)$
 - What if we want to model sequence length of, say, $L > 10^4$

Speed up Transformers


- Quadratic generation cost
 - $O(L^2)$ for length L : sequential generation $O(L) \times$ attention $O(L)$
- Make **attention** faster/better
 - Large-scale training: transformer-XL; XL-net (Zhilin Yang, et al, Google, 2020)
 - Projection tricks: Linformer (Facebook AI, $O(n)$ computation, 2020)
 - Math tricks: Performer (Google, $O(n)$ computation, 2020)
 - Sparse interactions:
 - Big Bird (Google, 2020), Multi-head Latent Attention (DeepSeek, 2024) <https://planetbanatt.net/articles/mla.html>
 - Fast and memory-efficient attention:
 - Flash Attention (Tri Dao, et al, 2022) and Ring Attention (Hao Liu, et al, 2023)
 - System engines for fast generation: vLLM (Berkeley) and SGLang (xAI & UCLA)
 - Even Parallel/Contextualized RNN (make RNN great again):
 - RWKV RNN (Open-Source, 2023) & Mamba (Gu, Albert and Tri Dao, 2023)
 - Reduce attention flatten issue when length grows: Scalable-Softmax (2025)

Speed up Transformers

- Quadratic generation cost
 - $O(L^2)$ for length L : sequential generation $O(L) \times$ attention $O(L)$
- Remark:
 - Ideally, attention can be computed in parallel given unbounded computation and memory bandwidth
- Can we accelerate autoregressive generation?
 - This is the key bottleneck for language model generation, which cannot be accelerated by hardware improvement


Beyond Autoregressive Generation

- WaveNet (Recap)
 - Let's use the language model notations
 - Output $y_1 \dots y_i$
 - Input $x_1 \dots x_i$ ($x_i = y_{i-1}$)
 - $p(y) = \prod_i p(y_i|x_1 \dots i)$
 - y_i can only be computed after y_{i-1}
 - $p(y_i) = N(\mu_\theta(x_1 \dots i), \exp^2(\alpha_\theta(x_1 \dots i)))$
 - $x_i \leftarrow y_{i-1}$
 - y_{i-1} is part of input of y_i
 - **Can we compute y_i without waiting?**




Beyond Autoregressive Generation

- WaveNet (Recap)
 - Let's use the language model notations
 - Output $y_1 \dots y_i$
 - Input $x_1 \dots x_i$ ($x_i = y_{i-1}$)
 - $p(y) = \prod_i p(y_i|x_1 \dots i)$
 - Reparameterization trick
 - $z_i \sim N(0,1)$
 - $y_i \leftarrow \mu_\theta(x_1 \dots i) + z_i \cdot \exp(\alpha_\theta(x_1 \dots i))$
 - $x_i \leftarrow y_{i-1}$




Beyond Autoregressive Generation

- WaveNet (Recap)
 - Let's use the language model notations
 - Output $y_1 \dots y_i$
 - Input $x_1 \dots x_i$ ($x_i = y_{i-1}$)
 - $p(y) = \prod_i p(y_i|x_1 \dots i)$
 - Reparameterization trick
 - $z_i \sim N(0,1)$
 - $y_i \leftarrow \mu_\theta(x_{1 \dots i}) + z_i \cdot \exp(\alpha_\theta(x_{1 \dots i}))$
 - $x_i \leftarrow y_{i-1}$
- x_i can be written as a function of $z_{1 \dots i}$ without compromising the representation power!
 - Each x is corresponding to a unique z ! (your homework)




Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)
 - Sequential modeling notation (ignore y)
 - Sequence tokens x & latent variable z
 - $p(x) = \prod_i p(x_i|x_1...i)$
 - Reparameterization trick
 - $z_i \sim N(0,1)$
 - $x_i \leftarrow \mu_\theta(z_1...i-1) + z_i \cdot \exp(\alpha_\theta(z_1...i-1))$
 - Parallel generation
 - First generate z and then x



Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)
 - Sequential modeling notation (ignore y)
 - Sequence tokens x & latent variable z
 - $p(x) = \prod_i p(x_i|x_1...i)$
 - Reparameterization trick
 - $z_i \sim N(0,1)$
 - $x_i \leftarrow \mu_\theta(z_1...i-1) + z_i \cdot \exp(\alpha_\theta(z_1...i-1))$
 - Parallel generation
 - First generate z and then x
 - **What about inference?**
 - Given x , how to compute $p(x)$ for MLE training?



Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)

- Sequential modeling notation (ignore y)

- Sequence tokens x & latent variable z

- $p(x) = \prod_i p(x_i | x_{1..i})$

- Reparameterization trick

- $z_i \sim N(0,1)$

- $x_i \leftarrow \mu_\theta(z_{1..i-1}) + z_i \cdot \exp(\alpha_\theta(z_{1..i-1}))$

- Parallel generation


- First generate z and then x

- Sequential inference $O(L)$

- x_i is a function of $z_1 \dots z_i$


- $z_i \leftarrow (x_i - \mu_i) / \exp(\alpha_i)$

- z_i can only be recovered after $z_1 \dots z_{i-1}$ is computed




Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)



Standard WaveNet

- Autoregressive generation $O(L \log L)$
- Parallel inference $O(\log L)$




Parallel WaveNet

- Parallel generation $O(\log L)$
- Autoregressive inference $O(L \log^8 L)$

Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)



Standard WaveNet

- Autoregressive generation $O(L \log L)$
- Parallel inference $O(\log L)$

Parallel WaveNet

- Parallel generation $O(\log L)$
- Autoregressive inference $O(L \log^{88} L)$

Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)
 - Key facts
 - Standard WaveNet is fast for training (inference is fast)
 - Parallel WaveNet is fast for serving (generation is fast)
 - Distillation by teacher-student framework!
 - Teacher: $p_T(x_i|x_{<i})$ a standard WaveNet for training on massive data
 - Student: $p_S(x_i|z_{<i})$ a parallel WaveNet for serving
 - p_S is trained by distillation from p_T
 - i.e., minimize the KL-difference between $p_S(x)$ and $p_T(x)$
 - Algorithm Sketch
 - Step 1: Train teacher $p_T(x_i|x_{<i})$ network and fix it
 - Step 2: Minimize the KL-difference $KL(p_S||p_T)$
 - Finally we use $p_S(x)$ for fast sampling

Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)
 - Key facts
 - Standard WaveNet is fast for training (inference is fast)
 - Parallel WaveNet is fast for serving (generation is fast)
 - Distillation by teacher-student framework!
 - Teacher: $p_T(x_i|x_{<i})$ a standard WaveNet for training on massive data
 - Student: $p_S(x_i|z_{<i})$ a parallel WaveNet for serving
 - p_S is trained by distillation from p_T
 - i.e., minimize the KL-difference between $p_S(x)$ and $p_T(x)$
 - Algorithm Sketch
 - Step 1: Train teacher $p_T(x_i|x_{<i})$ network and fix it
 - Step 2: Minimize the **KL-difference** $KL(p_S||p_T)$
 - Finally we use $p_S(x)$ for fast sampling


Pay attention to the order!

Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)
 - Teacher-Student Learning
 - Pretrain $p_T(x)$ and then train $p_S(x)$ by imitation learning
 - Distance measure for two distributions
 - KL divergence: $KL(p||q) = \mathbb{E}_{x \sim p} \left[\log \frac{p(x)}{q(x)} \right]$
 - Distillation (Imitation learning)
$$L(\theta) = KL(p_S||p_T) = \mathbb{E}_{x \sim p_S} [\log p_S(x; \theta) - \log p_T(x)]$$
 - Monte Carlo estimates for the expectation
 - **Key: sample from the student network!**
 - Sample $z \sim N(0, I)$, generate $x \sim p_S(x|z)$ (parallel)
 - Evaluate $p_S(x|z)$ (parallel since z is known)
 - Evaluate $p_T(x)$ (parallel since $p_T(x)$ is a standard autoregressive model)


Beyond Autoregressive Generation

- Parallel WaveNet (DeepMind, ICML 2018)
 - Teacher-Student Learning
 - Pretrain $p_T(x)$ and then train $p_S(x)$ by imitation learning
 - Speedup
 - 20x faster than real-time
 - 1000x faster than WaveNet
 - Google production
 - Remark
 - A reparameterization trick is assumed!
- What about language model?
 - Output are discrete tokens
 - No parameterization available



Beyond Autoregressive Generation


- Diffusion-based language model
 - Key idea: use a diffusion model to generate all tokens at once!
 - Idea#1: treat embeddings of tokens as images



- Pros: we can directly apply all techniques from diffusion models
- Cons: **length issue**, extremely high dimensions, poor generation quality

Beyond Autoregressive Generation



- Diffusion-based language model
 - Key idea: use a diffusion model to generate all tokens at once!
 - Idea#1: treat embeddings of tokens as images
 - Idea#2: define the denoising process over token masks



- Pros: reduce denoising operator to standard token prediction, reasonable quality
- Cons: **length issue**, no denoising acceleration anymore
- **Any simpler method for joint token predictions?**

Beyond Autoregressive Generation

- Multi-Token Prediction (ICML 2024)
 - Key insight $p(x_{1..L}) \approx \prod_i p(x_i)$
 - When the sequence length L is really short, we can break the sequential dependency
 - Predict K tokens jointly in parallel
 - $p(x_{i-K+1} \dots | x_{1..i-K}) \approx \prod_{j=1}^K p(x_{i-j+1} | x_{1..i-K})$
 - Small K & larger models higher gains



Beyond Autoregressive Generation

- Multi-Token Prediction (ICML 2024)
 - Key insight $p(x_{1..L}) \approx \prod_i p(x_i)$
 - When the sequence length L is really short, we can break the sequential dependency
 - Predict K tokens jointly in parallel
 - $p(x_{i-K+1} \dots i) \approx \prod_{j=1}^K p(x_{i-j+1} | x_{1..i-K})$
 - Understanding multi-token prediction
 - Using a **faster but worse model** (independent model) to approximate the **target distribution** (full language model)
 - **Other choice of fast sampling model?**
 - E.g., an LSTM, or a just smaller model

Beyond Autoregressive Generation

- Speculative Decoding (ICML 2023)
 - Key facts
 - A general language model $p(x)$ is fast at evaluation but slow at generation
 - A sampling model $q(x)$ is fast at generation but at low quality
 - Goal: adaptively use $q(x)$ to generate at **easy cases**
 - How to define **easy** cases?

Beyond Autoregressive Generation

- Speculative Decoding (ICML 2023)
 - Key facts
 - A general language model $p(x)$ is fast at evaluation but slow at generation
 - A sampling model $q(x)$ is fast at generation but at low quality
 - Goal: adaptively use $q(x)$ to generate at **easy cases**
 - MCMC Sampling!
 - We treat $q(x)$ as a proposal distribution
 - Given a partial prefix $x_{1\dots i}$, run q to generate next K tokens x'
 - Evaluate $p(x'|x)$ and $q(x'|x)$, accept x' with prob. $\min\left(1, \frac{p(x'|x)}{q(x'|x)}\right)$ (parallel)
 - If rejection, re-sample $x' \propto \max(0, p(x'|x) - q(x'|x))$ (autoregressive)
 - In practice, we can run speculative sampling for multiple K

Beyond Autoregressive Generation

[INST]Write a poem for my three year old[/INST]

[INST]Write a poem for my three year old[/INST]

Faster Transformer Generation

- Reparameterization and distillation
 - Pros: fast training and inference
 - Cons: only works for continuous values
- Diffusion-based language model
 - High complexity for generation and still low generation quality
- Multi-token prediction
 - Trade generation quality for speed
- Speculative decoding
 - Most general approach to speed up generation
 - Can be applied to any trained language model without modification

Summary

- Language Model & Sequence to Sequence Model
 - Fundamental ideas and methods for sequence modeling/tasks
- Attention Mechanism
 - So far the most successful idea for sequence data in deep learning
 - A scale/order-invariant representation
 - Transformer: a fully attention-based model for sequence data
- Speedup Transformers
 - Generation is the key bottleneck for transformer models
 - Acceleration by faster attention
 - Acceleration by non-autoregressive generation
 - Model changes: distillation, diffusion, multi-token prediction
 - Sampling methods: speculative decoding

Thanks

Deep Learning, Spring 2025
IIIS, Tsinghua University